



GAMECENTRIC

Alberto Barbati

alberto@gamecentric.com

# C++20 Text Formatting

Italian C++ Conference 2020

June 13<sup>th</sup>, Rome Internet



[www.italiancpp.org](http://www.italiancpp.org)



## Who am I?

C++ programmer since 1995

Works in videogame industry  
since 2000

Trainer of Game Programming at Digital  
Bros Game Academy

Follows the works of the C++ Committee  
since 2008





## What will we see today?

- C++20 introduces the new `<format>` header that provides facilities for general purpose text formatting
  - A replacement for `sprintf/stringstream`
  - Type-safe
  - Extensible
  - Localization friendly
  - Faster



**Nice, but... why?**



## sprintf and friends

- Well known, in the C library since the beginning
- General purpose
- Use a string with placeholders as a template

```
sprintf(buf, "The answer is %d", 42);
```

- Formats to a memory buffer, although interface is unsafe (possible buffer overrun). A "safe" version `sprintf_s` is available



## Issue: type-unsafe, not extensible

- Arguments are passed as `va_list` using ellipsis
- Placeholders contain information about the type
- Bad things may happen if the type implied by the placeholder does not match the type of the argument (this is partly mitigated by smart compilers/static analysis)
- No support for arguments of user-defined types



## Issue: mandatory global locale

- Number formatting is always locale-dependent and relies on global state
- Can be a nuisance if you need both "C" locale and user-locale output in the same program
- Can be a problem if program uses multi-threading

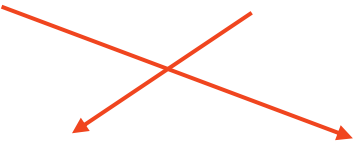


## Issue: argument matching

- Matching between placeholders and arguments is always positional

```
printf("Today is %s %d", month, dayOfMonth);  
// Today is June 13
```

```
printf("Oggi è il %d %s", dayOfMonth, month);  
// Oggi è il 13 giugno
```



- Localization may not be data-driven





## C++ iostreams

- Well known, in the C++ library since the beginning
- General purpose
- Use chains of operator <<

```
stream << "The answer is " << answer << "\\n";
```



## Advantage: type-safe and extensible

- Type-safe: The correct formatter function is selected at compile time, inferred from the argument type
- Extensible: it's easy to write formatters for user-defined types



## Issue: code-driven

- No placeholders, data is interspersed with context text strings

```
stream << "Today is " << month << " " << dayOfMonth;
```

- Complex formatting is not easily localizable
- Code is often verbose and difficult to read



## Issue: mandatory locale

- Formatting of numbers is always locale dependent
  - Relies on a per-stream locale object, defaulting to `std::locale::global()`, so use of global state is reduced
  - You can imbue a stream with a different locale object at any time
  - You can have both "C" locale and user-locale formatting, but it's cumbersome
  - There's an overhead cost even if you use only the "C" locale



## Issue: slow and bloated

- iostreams implementations are known for less than stellar performances, large code size and memory consumption
- The design doesn't allow much improvement, since it heavily relies on polymorphism in order to handle formatting, locale and buffering at the same time



## Issue: mandatory `std::string`

- The only standard way to format to a memory buffer is to use a `stringstream`
- It has a safe interface, but forces you to use `std::string`, which may require a dynamic allocation
- You need an extra copy to fill a pre-allocated buffer



## Alternative formatting libraries

- Boost.Format
- Folly.Format
- Loki.SafeFormat
- FastFormat
- QT QString
- Etc.

And...

**{fmt}**




## {fmt}

- {fmt} is the reference library that lead to the <format> proposal
- An open source library maintained by Victor Zverovich
- Available at <https://fmt.dev>
- MIT License
- Portable (GCC 4.8, Clang 3.0 and MSVC 19.0)
- Latest version requires C++11, but there's also a version for C++98





# {fmt} speed benchmark (floating point formatting)



shorter is faster (source <https://github.com/fmtlib/fmt>)



## Differences between `<format>` and `{fmt}`

- Only a subset of `{fmt}` has been included in C++20
- More than enough for everyday use
- A few extra `{fmt}` features might be considered for C++23



# Overview

**`std::format("The answer is {}", 42);`**

- Uses format strings with {}-bracketed placeholders
- Clearly inspired by Python and other languages



## Type safety

- Formatting is determined by the static type of the argument

<code>std::format(...)</code>	output
<code>"An integer: {}", 42</code>	An integer: 42
<code>"An ulonglong: {}", 42ull</code>	An ulonglong: 42
<code>"An fp number: {}", 1.234</code>	An fp number: 1.234
<code>"A string: {}", "hello"</code>	A string: hello



## Arguments matching

- By default, placeholders are matched with arguments positionally

std::format(...)	output
"{}", {}, {}, 42, 1.234, "hello"	42, 1.234, hello





## Arguments matching

- The default *automatic* indexing can be overridden by manually specifying indices

<code>std::format(...)</code>	output
<code>"Today is {0}, {1}", month, day</code>	Today is June, 13
<code>"Oggi è il {1} {0}" , month, day</code>	Oggi è il 13 Giugno



## Arguments matching

- Manual indices can be omitted and/or repeated

<code>std::format(...)</code>	output
<code>"Omission {0} {2}", 0, 1, 2</code>	Omission 0 2
<code>"Repetition {0} {0} {0}", 42</code>	Repetition 42 42 42



## Arguments matching

- You can't mix automatic and manual indexing in the same string

<code>std::format(...)</code>	output
"Looking for troubles {} {1}"	throws <code>std::format_error</code>





## Types supported by the library

- The library supports, out of the box:
  - Arithmetic types: all integers and floating-point types, plus `bool`, `char` and `wchar_t`, but excluding `char8_t`, `char16_t` and `char32_t`
  - String types: null-terminated (`const char*`), `char` arrays, `std::string` and `std::string_view`
  - Pointer types: `void*` and `nullptr`
  - Time/date types defined in `<chrono>`



## Format specifiers

- To customize formatting, you can provide optional format specifiers after the index, if present
- Format specifiers are always introduced by a colon

```
std::format("The answer is {:specifiers}", 42);
```

```
std::format("Today is {0:specifiers}, {1: specifiers}", month, day);
```



## Format specifiers

- The interpretation of the format specifiers is determined by the type of the matching argument
- If a type doesn't support a specific format specifier, a `std::format_error` exception is thrown
- The library design is open to the possibility to have the format string parsed and checked at compile time, but C++20 still lacks language support to have this performed automatically
- `{fmt}` implements compile time checks, but requires explicit opt-in



## Format specifiers for common types

- The general format specifier for the standard-supported types is:

***fill-and-align sign # 0 width precision L type***

- Every piece is optional
- A few pieces only make sense for specific types
- Refer to [format.string.std] for details

Also have a look to  
<https://fmt.dev/latest/syntax.html>



## Common specifiers

- A few pieces work as they do in `sprintf` (or similar):
  - *sign* controls the presence of a + sign before positive numbers
  - *#* switches to an "alternate" formatting
  - *0* produces 0-padding for numbers
  - *width* specifies the minimum output width
  - *precision* specifies the maximum output width (for strings) or controls the number of digits (for floating point numbers)



## Alignment: examples

<code>std::format(...)</code>	output
<code>" {:8} ", 42</code>	<code>       42 </code>
<code>" {:*&gt;8} ", 42</code>	<code> *****42 </code>
<code>" {:*&lt;8} ", 42</code>	<code> 42***** </code>
<code>" {:.*^8} ", 42</code>	<code> ***42*** </code>
<code>" {:8} ", "hello"</code>	<code> hello    </code>
<code>" {:*&gt;8} ", "hello"</code>	<code> ***hello </code>
<code>" {:*&lt;8} ", "hello"</code>	<code> hello*** </code>
<code>" {:.*^8} ", "hello"</code>	<code> *hello** </code>



## Opt-in locale support

- By default, numbers are formatted in the "C" locale, but you can opt-in for a localized formatting using the L specifier

<code>std::format(...)</code>	Output (in the IT_it locale)
<code>"{0}", 1234.56</code>	1235.56
<code>"{0:L}", 1234.56</code>	1234,56

- Formatting functions can either use the global locale or have the locale passed as an argument

The `fmt` library before 6.2.1  
uses `n` instead of `L`



## Integer formatting examples

<code>std::format(...)</code>	output
<code>"{:d}", 42</code>	42
<code>"{:b}", 42</code>	101010
<code>"{:o}", 42</code>	52
<code>"{:x}", 42</code>	2a
<code>"{:#x}", 42</code>	0x2a
<code>"{:08x}", 42</code>	0000002a
<code>"{:#08x}", 42</code>	0x00002a





## Floating point formatting examples

- Floating point formatting rely upon C++17 function `to_chars`:

<code>std::format(...)</code>	<code>to_chars</code> equivalent	output
<code>"{}", 3.141592654</code>	(shortest round-trip)	<code>3.141592654</code>
<code>"{:f}", 3.141592654</code>	<code>chars_format::fixed</code>	<code>3.141593</code>
<code>"{:g}", 3.141592654</code>	<code>chars_format::general</code>	<code>3.14159</code>
<code>"{:e}", 3.141592654</code>	<code>chars_format::scientific</code>	<code>3.141593e+00</code>
<code>"{:a}", 3.141592654</code>	<code>chars_format::hex</code>	<code>0x1.921fb54524550p+1</code>



## Formatting functions

- The simplest formatting function is `std::format`, which returns the formatted string as a `std::string` object:

```
template <class Args...>  
string format(string_view fmt, const Args&... args);
```



## Example

```
string s = format("The answer is {}", 42);  
cout << format("Today is {}, {}\n", "June", 13);
```



## Formatting functions

- Actually we have four overloads of format:

```
string format(string_view fmt, const Args&... args);
```

```
string format(const locale& loc, string_view fmt, const Args&... args);
```

```
wstring format(wstring_view fmt, const Args&... args);
```

```
wstring format(const locale& loc, wstring_view fmt, const Args&... args);
```

- Similarly for all other formatting functions. For the sake of brevity, I will omit the overloads for wstring and with the explicit locale object



## Formatting into containers

- You can format directly into a container, without allocating a string:

```
template<class Out, class... Args>  
Out format_to(Out out,  
              string_view fmt,  
              const Args&... args);
```

- This function takes an output iterator and returns an iterator past the end of the formatted text
- No null termination is appended



## Example

```
std::vector<char> v;  
format_to(back_inserter(v), "The answer is {}", 42);  
  
format_to(ostreambuf_iterator(cout),  
          "Today is {}, {}\\n", "June", 13);
```



## Formatting into fixed size buffers

- To make it safer to work with pre-allocated buffers of fixed size, there's a variant that takes an explicit maximum size:

```
template <class Out, class... Args>
format_to_n_result<Out>
format_to_n(Out out,
            iter_difference_t<Out> n,
            string_view fmt,
            const Args&... args);
```



## Formatting into fixed size buffers

- The return type contains both the iterator and the full size of the formatted string

```
template <class Out>
struct format_to_n_result
{
    Out out;
    iter_difference_t<Out> size;
};
```





## Example (buffer big enough)

```
char buffer[20];  
auto [out, size] = format_to_n(buffer, sizeof(buffer),  
                                "The answer is {}", 42);  
  
// buffer = "The answer is 42□□□□"  
// out == buffer + 16  
// size == 16
```

uninitialized



## Example (buffer too small)

```
char buffer[10];  
auto [out, size] = format_to_n(buffer, sizeof(buffer),  
                                "The answer is {}", 42);
```

```
// buffer = "The answer"  
// out == buffer + 10  
// size == 16
```

no null-termination  
past-the-output iterator  
required buffer size



## Formatted size

- If you just need to know the size of the formatted string without actually formatting it, you can use

```
template<class... Args>  
size_t formatted_size(string_view fmt,  
                      const Args&... args);
```



## Fighting code bloat

- Since all format functions we have seen so far are templates, it's important to avoid code bloat
- The library actively fights that by:
  - Packing all the variadic arguments in a type-erased struct with the function `make_format_args()`
  - Doing the actual formatting in a non-template function

```
string vformat(string_view fmt, format_args args);
```



## Example of {fmt} generated code

```
#include <fmt/core.h>
```

```
int main()
{
    fmt::print(
        "The answer is {}. ", 42);
}
```

```
main: # @main
    sub rsp, 24
    mov qword ptr [rsp], 42
    mov rcx, rsp
    mov edi, offset .L.str
    mov esi, 17
    mov edx, 2
    call fmt::v5::vprint(...)
    xor eax, eax
    add rsp, 24
    ret
.L.str:
    .asciz "The answer is {}."
```



## Custom formatting example: the type

```
struct Conference
{
    std::string title;
    std::string location;
    std::chrono::sys_days date;
};
```

```
Conference italianCpp20 { "Italian C++ Conference",
                          "Internet", // previously Rome
                          2020y/June/13 };
```



## Custom formatting example: desired output

```
format("Event: {}", italianCpp20);
```

```
// output:
```

```
// Event: Italian C++ Conference - 2020-06-13, Internet
```



## Custom formatters

- Formatting can be extended to support user-defined types by specializing the `std::formatter` template
- The formatter needs to implement both these steps:
  - Parsing the format specifier
  - Do the formatting





## Custom formatting example: the plan

```
namespace std
{
    template<>
    struct formatter<Conference>
    {
        template <class ParseCtx>
        constexpr typename ParseCtx::iterator parse(ParseCtx& ctx)
        {
            // TODO: parsing
        }

        template <class FormatCtx>
        typename FormatCtx::iterator format(const Conference& c, FormatCtx& ctx)
        {
            // TODO: formatting
        }
    };
}
```



## Custom formatting example: the plan (C++20 syntax)

```
template<>
struct std::formatter<Conference>
{
    template <class ParseCtx>
    constexpr ParseCtx::iterator parse(ParseCtx& ctx)
    {
        // TODO: parsing
    }

    template <class FormatCtx>
    FormatCtx::iterator format(const Conference& c, FormatCtx& ctx)
    {
        // TODO: formatting
    }
};
```



## Custom formatting example: parse part

```
template <class ParseCtx>
constexpr ParseCtx::iterator parse(ParseCtx& ctx)
{
    // parse the range [ctx.begin(), ctx.end())
    // return the iterator to the first unparsed char
    // throw std::format_error on error
}
```



## Custom formatting example: the range

"Lorem { :specifiers } ipsum"



ctx.begin()



ctx.end()



## Custom formatting example: parse part (no specifiers)

```
template <class ParseCtx>
constexpr ParseCtx::iterator parse(ParseCtx& ctx)
{
    auto it = ctx.begin();
    if (it == ctx.end() || *it != '}')
        throw std::format_error();
    return it;
}
```



## Custom formatting example: format part

```
template <class FormatCtx>
FormatCtx::iterator format(const Conference& c, FormatCtx& ctx)
{
    // format text to ctx.out()
    // return iterator past the formatted text
}
```



## Custom formatting example: format part

```
template <class FormatCtx>
FormatCtx::iterator format(const Conference& c, FormatCtx& ctx)
{
    return format_to(ctx.out(),
                     "{} - {:%F}, {}",
                     c.title, c.date, c.location);
}
```

%F formats a date in ISO format  
YYYY-MM-DD



## A slightly more complex example

```
format("Event: {}", italianCpp20);  
// Event: Italian C++ Conference – 2020-06-13, Internet
```

```
format("Event: {:L}", italianCpp20);  
// Event: Italian C++ Conference – June 13 2020, Internet
```

```
format(locale("IT_it"), "Event: {:L}", italianCpp20);  
// Event: Italian C++ Conference – 13 Giugno 2020, Internet
```





## A slightly more complex example

```
template<>
struct std::formatter<Conference>
{
    bool localizedFormat = false;

    template <class ParseCtx>
    constexpr ParseCtx::iterator parse(ParseCtx& ctx);

    template <class FormatCtx>
    FormatCtx::iterator format(const Conference& c, FormatCtx& ctx);
};
```



## A slightly more complex example: parsing

```
template <class ParseCtx>
constexpr ParseCtx::iterator parse(ParseCtx& ctx)
{
    auto it = ctx.begin();
    if (it != ctx.end() && *it == 'L')
    {
        localizedFormat = true;
        ++it;
    }
    if (it == ctx.end() || *it != '}')
        throw std::format_error();
    return it;
}
```



## A slightly more complex example: formatting

```
template <class FormatCtx>
FormatCtx::iterator format(const Conference& c, FormatCtx& ctx)
{
    if (localizedFormat)
        return format_to(ctx.locale(),
                        ctx.out(),
                        "{} - {:%x}, {}",
                        c.title, c.date, c.location);
    else
        // as before
}
```

%x produces the localized date format



## Another example: enumerations

```
enum class Color { Red, Green, Blue };  
  
format("Color: {}", Color::Red);  
// Color: Red
```



## Another example: enumerations

```
const char* colors[] = { "Red", "Green", "Blue" };

template<>
struct std::formatter<Color> : std::formatter<const char*>
{
    // parse is inherited from base class

    auto format(Color c, auto& ctx)
    {
        // format is delegated to base class
        return formatter<const char*>::format(colors[(int)c], ctx);
    }
};
```



## Recap

- The C++20 `<format>` header provides facilities for general purpose text formatting, addressing major issues of `sprintf/stringstream`:
  - Type-safe
  - Extensible
  - Localization friendly
  - Performances and code bloat
- You can try it right now by using the `{fmt}` library



**Thanks for your attention**



# Questions?

Alberto Barbati  
alberto@gamecentric.com  
Twitter @gamecentric