

# The Silicon Valley coding interview

---

Nicolo Valigi    `nicolovaligi.com`

June 13, 2020

# Silicon Valley

The Silicon Valley tech ecosystem is huge, and has its own peculiar way of interviewing people for jobs.



# Is the coding interview too hard? Or just wrong?

Not everybody is excited about being tested for coding before getting a job..



**Max Howell** @mxcl · Jun 10, 2015



Google: 90% of our engineers use the software you wrote (Homebrew), but you can't invert a binary tree on a whiteboard so fuck off.

💬 516

↻ 7.4K

❤ 12.1K



A common complaint is that the typical coding interview is not representative of the actual day-to-day job.

# Goals for the talk

## Who is this talk for?

- You're thinking about applying to one of the tech companies
- You want to laugh about how ridiculous people are
- You just like code puzzles and algorithms

## This talk will not:

- replace a degree in CS
- land you a job at Google tomorrow
- teach you anything about C++

# Plan for the talk

- Recruiting process and timeline
- What interview questions look like
- Families of algorithmic questions
- Tips for preparation

## Example question

Some examples of questions I actually got during interviews:

*Given an array of numbers and an integer  $n$ , find the number of triplets  $(a, b, c)$  such that  $a+b+c$  is less than  $n$ .*

or:

*Implement a Tic-Tac-Toe game.*

# The recruiting process

- You call up a recruiter (or a recruiter calls you)
- Short call with the recruiter (non-technical)
- 1-hour programming test remotely (eg [coderpad.io](https://coderpad.io))
- 5 1-hour programming tests on-site
- Feedback and maybe a job offer

# The technical interview

---



# What questions to expect

- Most interview questions are about **algorithms** and **data structures**.
- Explained in just a few minutes in very abstract terms.
- Little domain knowledge needed for most of them.

# Families of questions

Most questions will fall into one of these categories:

- *linear things*: arrays, (linked) lists, searching, and sorting
- *stacks and queues*
- *tree-ish things*: binary search trees, heaps
- *graphs*: search, connected components
- *dynamic programming and recursion*
- *system design*

# Array problems

---

## Arrays (1/)

*Given  $\text{int}[] A$  and  $\text{int } N$ , find two indices  $i$  and  $j$  in  $A$  such that  $A[i] + A[j] == N$ . Assume that the solution exists and is unique. Each element of  $A$  can only be used once.*

*Example:*

`TwoSum({2, 7, 11, 15}, 9) -> {0, 1}`

## Arrays (2/)

A trivial solution uses a nested for loop:

```
for (int i = 0; i < A.size(); i++) {  
    for (int j = i+1; j < A.size(); j++) {  
        if (A[i] + A[j] == N) {  
            return {i, j};  
        }  
    }  
}
```

Works but has poor **asymptotic complexity** (ie, it becomes slow with large arrays).

## Intermission: asymptotic complexity

Measures how computational cost (or memory) grows with increasing input size.

- **constant** complexity: same amount of CPU/memory no matter the size of the input.  
*Example:* get the first element of the array.
- **linear** complexity: grows linearly with the input size. *Example:* sum all the elements of the array.
- **quadratic** complexity
- **exponential** complexity: usually, this means you have to do better..

The interviewer will often ask you what the a.c. of your solution is, and maybe to find one with lower a.c.

## Arrays (3/)

- What's the time and space complexity of the trivial solution?
- Can we write a better solution?

## Arrays (3/)

- What's the time and space complexity of the trivial solution?
- Can we write a better solution?

We can go from quadratic to linear time complexity if we take advantage of extra memory:

```
std::unordered_map<int, size_t> valueToIndex;
for (int i = 0; i < A.size(); i++) {
    const int rem = N - A[i];
    if (auto it = valueToIndex.find(rem); it != valueToIndex.end()) {
        return {it->second, i};
    }
    valueToIndex[A[i]] = i;
}
```



## Stack problems

---

# Stacks (1/)

Stacks are the prototypical *first-in, first-out* datastructure.

*Different UNIX pathnames can represent the same directory on disk, Write a function that returns the shortest possible equivalent pathname to the input.*

*Example:*

`/home/niko//code/../../code/./project -> /home/niko/code/project`

## Stacks (2/)

The idea is to push each path fragment to a stack, and pop on “..”.

## Stacks (2/)

The idea is to push each path fragment to a stack, and pop on “..”.

```
stack<string> fragments;

stringstream ss(path);
string fragment;

while (getline(ss, fragment, '/') {
    if (fragment == "." || fragment == "") { }
        continue;
    } else if (fragment == "..") {
        } else if (fragment == "..") {
            if (fragments.empty()
                || fragments.back() == "..") {
                fragments.push(fragment);
            } else {
                fragments.pop();
            }
        } else {
            fragments.push(fragment);
        }
    }
```

## Stacks (3/)

And finally we assemble the result back together from the stack:

```
string result;  
while (!fragments.empty()) {  
    result += fragments.top();  
    fragments.pop();  
}  
return result;
```

- Stacks are a good choice for problems that can be solved incrementally.

# Tree problems

---

# Trees (1/)

A (binary) tree is a hierarchical data structure, where each *node* holds references to a left and right children.


```
template <typename T>
struct Node {
    T value;
    Node* left;
    Node* right;
};
```

Many operations on trees are naturally recursive.

## Trees (2/)

*Determine if a binary tree is symmetric around its center.*

This tree is symmetric:



This tree is *not* symmetric:





## Trees (3/)

```
bool IsSymmetric(Node* tree) {  
    return IsMirror(tree->left, tree->right);  
}
```

## Trees (3/)

```
bool IsSymmetric(Node* tree) {  
    return IsMirror(tree->left, tree->right);  
}  
  
bool IsMirror(Node* subtree1, Node* subtree2) {  
    if (subtree1 == nullptr && subtree2 == nullptr) {  
        return true;  
    }
```

## Trees (3/)

```
bool IsSymmetric(Node* tree) {  
    return IsMirror(tree->left, tree->right);  
}  
  
bool IsMirror(Node* subtree1, Node* subtree2) {  
    if (subtree1 == nullptr && subtree2 == nullptr) {  
        return true;  
    } else if (subtree1 != nullptr && subtree2 != nullptr) {  
        return subtree1->value == subtree2->value &&
```

## Trees (3/)

```
bool IsSymmetric(Node* tree) {  
    return IsMirror(tree->left, tree->right);  
}  
  
bool IsMirror(Node* subtree1, Node* subtree2) {  
    if (subtree1 == nullptr && subtree2 == nullptr) {  
        return true;  
    } else if (subtree1 != nullptr && subtree2 != nullptr) {  
        return subtree1->value == subtree2->value &&  
            IsMirror(subtree1->left, subtree2->right) &&  
            IsMirror(subtree1->right, subtree2->left);  
    }  
}
```

## Trees (3/)

```
bool IsSymmetric(Node* tree) {  
    return IsMirror(tree->left, tree->right);  
}  
  
bool IsMirror(Node* subtree1, Node* subtree2) {  
    if (subtree1 == nullptr && subtree2 == nullptr) {  
        return true;  
    } else if (subtree1 != nullptr && subtree2 != nullptr) {  
        return subtree1->value == subtree2->value &&  
            IsMirror(subtree1->left, subtree2->right) &&  
            IsMirror(subtree1->right, subtree2->left);  
    } else {  
        return false;  
    }  
}
```

## Graph problems

---

## Graph search (1/)

*You're given two words (`beginWord` and `endWord`) and a list of words `wordList`. Find the length of a sequence of transformations from `beginWord` to `endWord` where only one letter can be changed at a time, and all words must be in `wordList`. Return 0 if no such sequence exists.*

```
findSequence("hit", "cog", {"hot", "dot", "lot", "log", "cog"})  
-> 5
```

Example transformation sequence:

```
"hit" -> "hot" -> "dot" -> "dog" -> "cog"
```

## Graph search (2/)

```
queue<string> q;  
q.push(beginWord);  
int seqLen = 1;
```



## Graph search (2/)

```
queue<string> q;
q.push(beginWord);
int seqLen = 1;

while (!q.empty()) {
    const int n = q.size();
    for (int i=0; i < n; i++) {
        const string cur = q.front();
        q.pop();
        dictionary.erase(word); // Don't use this word again.
        if (cur == endWord) { return seqLen; } // Done.
        // List neighbors and push them into q... <=====
    }
    seqLen++;
}
```

## Graph search (3/)

Here's how we generate new candidates for exploration. Remember `cur` is the current word we're looking at.

```
for (int j = 0; j < cur.size(); j++) {  
    char c = cur[j];  
    for (int k = 0; k < 26; k++) {  
        string cand = cur;  
        // Overwrite the j-th character  
        cand[j] = 'a' + k;  
        if (dictionary.find(cand) != dictionary.end()) {  
            q.push(cand);  
        }  
    }  
}
```

# Dynamic Programming problems

---

# Dynamic Programming (1/)

*You're given a list of coin denominations `values` and an integer `amount`. Find the minimum number of coins needed to add up to that amount. If no combination of `values` can add up to `amount`, return `-1`.*

Example:

`NumCoinsRequired(11, {1, 2, 5}) -> 3`

made up by  $5 + 5 + 1 = 11$

## Dynamic Programming (2/)

This problem satisfies the *optimal substructure* property: the optimal solution  $X(\text{amount})$  satisfies the following relation:

$$X(\text{amount}) = X(\text{amount} - \text{coin}) + 1$$

where `coin` is the value of *some* coin, but we don't know which yet. So we try all of them and pick the minimum:

$$X(\text{amount}) = \min_{i=0..n-1} X(\text{amount} - c_i) + 1$$

This allows us to start from small values of `amount` and reuse lots of computation as we move towards bigger values.

## Dynamic Programming (3/)

```
vector<int> dp(amount+1, amount+1);  
dp[0] = 0;  
  
for (int i = 1; i <= amount; i++) {  
    for (int j = 0; j < coins.size(); j++) {  
        if (coins[j] <= i) {  
            dp[i] = min(dp[i], dp[i - coins[j]] + 1);  
        }  
    }  
}  
  
return dp[amount] > amount ? -1 : dp[amount];
```

# System design problems

---

*You have to sort more data that fits in the memory of any single machine. What do you suggest to get the job done quickly?*

- Open-ended question with several potential solutions and tradeoffs to explore.
- Higher level topics than pure datastructures (eg networking, computer architecture, etc..)



## System design (2/)

So, what do we do?

So, what do we do?

- Spill data to disk (order of magnitudes slower than memory or network).

## System design (2/)

So, what do we do?

- Spill data to disk (order of magnitudes slower than memory or network).
- Distribute the data to multiple machines, but how?

## System design (2/)

So, what do we do?

- Spill data to disk (order of magnitudes slower than memory or network).
- Distribute the data to multiple machines, but how?
- At first glance, *merge sort* might be a good algorithm (just merge the data to/from multiple machines at the same time).

## System design (2/)

So, what do we do?

- Spill data to disk (order of magnitudes slower than memory or network).
- Distribute the data to multiple machines, but how?
- At first glance, *merge sort* might be a good algorithm (just merge the data to/from multiple machines at the same time).
- If the data is segmented by a natural key, use *bucketing* techniques to spread the load across multiple machines.

## Wrapping up

---

# Recap of problems

- *arrays*: usually the goal is to minimize time/space complexity.
- *stacks*: applies to problems with an incremental/sequential structure.
- *trees*: the trick is to identify a common property across levels and write it down recursively.
- *graphs*: the trick is to write down a problem in terms of edges, nodes, and connectivity between the two.
- *Dynamic Programming*: the trick is to find a repeating structure of smaller problems.
- *System Design*: good opportunity to show off your real-world experience.

## Things we left out:

- (Linked) lists and queues.
- Backtracking with recursion.
- Searching and sorting.

## Wrapping up

- Algorithms are fun. Not so much when you're under pressure in front of a whiteboard.
- The interviewer wants you to do well (so they'll have one more team member to help out).
- Start with simple solutions, then iterate to improve the time/space complexity.
- Some interviewers give lots of tips if you're stuck, some are dead silent. Asking helps.
- Practice is everything.



# References

- *Aziz, Lee, Prakash*, Elements of Programming Interviews: The Insiders' Guide
- EPI judge [github.com](https://github.com/donnemartin/epi-judge)
- [leetcode.com](https://leetcode.com)
- *Steven S. Skiena*, The Algorithm Design Manual

**Thanks**

---